
Homework 2: Linear Classification

There is a mathematical component and a programming component to this homework. Please submit your PDF to Canvas, and push everything in Github.

This homework is about multi-class classification. Whereas in more simple classification models we build classifiers that discriminate between two classes, in multi-class regression, we discriminate between three or more classes. As usual, we imagine that we have the input matrix $X \in \mathbb{R}^{N \times D}$ (or perhaps they have been mapped to some basis Φ , without loss of generality) but that our outputs are now “one-hot coded”. What that means is that, if there are K output classes, rather than representing the output labels as integers $1, 2, \dots, K$, we represent them as a binary vectors of length K . These vectors are zero in each component except for the one corresponding to the correct label, and that entry has a one. So, if there are 7 classes and a particular datum has label 3, then the target vector would be $[0, 0, 1, 0, 0, 0, 0]$ (assuming the labels are 1-indexed).

In the first problem, you will be exploring the properties of the softmax function, which is central to multiclass logistic regression. In the second problem, we will have you dive into the matrix algebra and methods behind generative classifications. Finally, in the third problem, you will implement a generative classifier and logistic regression from close to scratch, and the first two problems should inform this!

Problem 1 (Properties of Softmax , 5pts)

Logistic regression is a discriminative probabilistic model: a prediction consists of a distribution over the different classes. In other words, logistic regression outputs a vector of nonnegative numbers that sum to one.

The softmax function generalizes the logistic sigmoid to the case of K classes: it takes as input a vector, and outputs a K dimensional vector in the range $[0, 1]$ whose components sum to 1:

$$\sigma(\mathbf{z}) = \text{softmax}(\mathbf{z}) = \frac{\exp(\mathbf{z})}{\sum_i \exp(z_i)}$$

In logistic regression, we often use the softmax-based parameterization over K vectors $\{\mathbf{w}_k\}$:

$$\Pr(t_{nk} = 1 \mid \mathbf{X}, \{\mathbf{w}_{k'}\}_{k'=1}^K) = \frac{\exp\{\mathbf{w}_k^\top \mathbf{x}_n\}}{\sum_{k'=1}^K \exp\{\mathbf{w}_{k'}^\top \mathbf{x}_n\}}.$$

Here we're using $t_{nk} = 1$ to indicate the probability that the n th entry is assigned to the k th class.

Softmax is a crucial function in logistic regression, and you will see it again in other models, such as neural networks. So, we want you to start gaining the intuitions for the properties of softmax, and for common methods that employ it.

Show that:

1. The output of the softmax function is always a vector with non-negative components that are at most 1.
2. The output of the softmax function forms a distribution (the components sum to 1).
3. Softmax preserves order. This means that if the elements of \mathbf{z} have some order, then the elements of $\sigma(\mathbf{z})$ have the same order.
4. Equation 4.106 from Bishop holds
5. Using your answer to the previous question, show that equation 4.109 holds. By the way, this may be useful for Problem 3!

Solution

1- $\forall x \in \mathbb{R}, \exp(x) > 0$ so $\forall j, \frac{\exp(z_j)}{\sum_i \exp(z_i)} > 0$ hence all components are non negative.

$\forall j, \exp(z_j) \leq \exp(z_j) + \sum_{i=1}^{j-1} \exp(z_i) + \sum_{i=j+1}^N \exp(z_i)$ since the sum are sums of strictly positive numbers.

hence $\forall j, \frac{\exp(z_j)}{\sum_i \exp(z_i)} < 1$.

$$2- \sum_j \frac{\exp(z_j)}{\sum_i \exp(z_i)} = \frac{1}{\sum_i \exp(z_i)} \sum_j \exp(z_j) = 1$$

3- if $z_j \leq z_{j+1}, \exp(z_j) \leq \exp(z_{j+1})$ since \exp is an increasing function.

$\frac{\exp(z_j)}{\sum_i \exp(z_i)} \leq \frac{\exp(z_{j+1})}{\sum_i \exp(z_i)}$ since we divided by a positive scalar, the sum of positive numbers.

4- First, let's look at the case where $j \neq k$

$$\frac{\partial y_k}{\partial a_j} = \exp(a_k) \cdot \frac{\partial}{\partial a_j} \frac{1}{\sum_i \exp(a_i)} = \exp(a_k) \left[\frac{-\exp(a_j)}{(\sum \exp(a_i))^2} \right]$$

$$\frac{\partial y_k}{\partial a_j} = \frac{\exp(a_k)}{\sum_i \exp(a_i)} \cdot \frac{-\exp(a_j)}{\sum_i \exp(a_i)} = y_k \cdot (-y_j) = y_k(I_{kj} - y_j) \text{ since when } j \neq k, I_{kj} = 0.$$

Now, $k = j$,

$$\frac{\partial y_j}{\partial a_j} = \frac{\partial}{\partial a_j} \exp(a_j) \cdot \frac{1}{\sum_i \exp(a_i)}$$

by using the product formula $(u.v)' = u'v + uv'$, we get:

$$\frac{\partial y_j}{\partial a_j} = \frac{\exp(a_j)}{\sum_i \exp(a_i)} + \exp(a_j) \cdot \frac{-\exp(a_j)}{\sum_i \exp(a_i)} = y_j(I_{jj} - y_j)$$

5-

$$E(w_1, \dots, w_K) = -\ln p(T|w_1 \dots w_K) = -\sum_{n=1}^N \sum_{k=1}^K t_{nk} \ln(y_{nk})$$

$$\nabla_{w_j} E(w_1, \dots, w_K) = -\sum_{n=1}^N \sum_{k=1}^K \nabla_{w_j} t_{nk} \ln(y_{nk}) = -\sum \sum t_{nk} \nabla_{w_j} \ln(y_{nk})$$

$$\nabla_{w_j} E(w_1, \dots, w_K) = -\sum \sum \frac{t_{nk}}{y_{nk}} \frac{\partial y_{nk}}{\partial w_j}$$

$$\nabla_{w_j} E(w_1, \dots, w_K) = -\sum \sum \frac{t_{nk}}{y_{nk}} \frac{\partial y_{nk}}{\partial a_j} \frac{\partial a_j}{\partial w_j} \text{ (change of variables)}$$

$$\text{we have } \frac{\partial a_j}{\partial w_j} = \phi_n$$

$$\nabla_{w_j} E(w_1, \dots, w_K) = -\sum \sum \frac{t_{nk}}{y_{nk}} \cdot y_{nk} (I_{kj} - y_{nj}) \phi_n \text{ (from previous question)}$$

$$\nabla_{w_j} E(w_1, \dots, w_K) = -\sum \sum t_{nk} (I_{kj} - y_{nj}) \phi_n$$

$$\nabla_{w_j} E(w_1, \dots, w_K) = -\sum_{n=1}^N \phi_n \sum_{k=1}^K (t_{nk} \cdot I_{kj} - t_{nk} \cdot y_{nj})$$

$$\nabla_{w_j} E(w_1, \dots, w_K) = -\sum_{n=1}^N \phi_n \left[\sum_{k=1}^K t_{nk} \cdot I_{kj} - \sum_{k=1}^K t_{nk} \cdot y_{nj} \right]$$

$$\nabla_{w_j} E(w_1, \dots, w_K) = \sum_{n=1}^N \phi_n \left[y_{nj} \sum_{k=1}^K t_{nk} - t_{nj} \right]$$

we have $\sum_{k=1}^K t_{nk} = 1$, hence:

$$\nabla_{w_j} E(w_1, \dots, w_K) = \sum_{n=1}^N \phi_n (y_{nj} - t_{nj})$$

Problem 2 (Moosooar matrix calculus , 10pts)

Note - this problem appears longer than it is, since we broke up one problem into separate parts rather than having you do all of these steps at once. Many of these subparts may be just one or two lines.

Consider a generative K-class model. We define the class prior with vector $\vec{\pi}$: $\mathbb{P}(\mathcal{C}_k) = \pi_k$. We define the class-conditional densities $\mathbb{P}(\phi|\mathcal{C}_k)$ where ϕ is the input feature vector. Consider the data set $\{\phi_n, \mathbf{t}_n\}$ where $n = 1 \dots N$ where $\mathbf{t}_n \in \{0, 1\}^K$ is a one-hot encoded target vector. This means that \mathbf{t}_n is 0 everywhere, except for in the k th position, where k is the class assigned to the n th feature vector.

1. Write out the complete-data log-likelihood of the data set using only the notations introduced in the problem formulation above.

$$\ln \mathbb{P}(\{\phi_n, \mathbf{t}_n\} | \{\pi_k\}) = ?$$

2. Since the prior forms a distribution, it has the constraint that $\sum_k \pi_k - 1 = 0$. Using the hint at the end of the exercise, give the expression for the maximum-likelihood estimator for the prior class-membership probabilities:

$$\hat{\pi}_k = ?$$

Make sure to write out the intermediary equation you need to solve to obtain this estimator. Double-check your answer: the final result should be very intuitive!

We will suppose for the remaining questions of this exercise that the class-conditional probabilities are given by gaussian distributions with the same covariance matrix:

$$\mathbb{P}(\phi|\mathcal{C}_k) = \mathcal{N}(\phi | \vec{\mu}_k, \Sigma)$$

3. Write out the gradient of log-likelihood with respect to vector μ_k . Write the expression in matrix form as a function of the variables defined throughout this exercise. Simplify as much as possible for full credit.
4. Write out the maximum-likelihood estimator for vector μ_k . Once again, your final answer should seem intuitive.
5. Write out the gradient for the log-likelihood with respect to the covariance matrix Σ . Even though the log-likelihood function is a scalar function, since you are differentiating with respect to a *matrix*, the resulting expression should be a matrix!
6. Express the maximum likelihood estimator of the covariance matrix.

Hint. When maximizing a function f with respect to an equality constraint that needs to be met at the optimum (which can always be written as $g(x) = 0$), we introduce a Lagrange multiplier λ and maximize:

$$\max_x f(x) + \lambda g(x)$$

Cookbook formulas. Here are some formulas you might want to consider using to compute difficult gradients. You can use them as is in the homework without proof. If you are looking to hone your matrix calculus skills, try to find different ways to prove these formulas yourself (will not be part of the evaluation of this homework). In general, you can use any formula from the matrix cookbook, as long as you cite it. We opt for the following common notation: $X^{-T} := (X^{-1})^{-T} = (X^T)^{-1}$

$$\frac{\partial a^T X^{-1} b}{\partial X} = -X^{-T} a b^T X^{-T}$$
$$\frac{\partial \ln |\det(X)|}{\partial X} = X^{-T}$$

Solution

1-

$$\ln \mathbb{P}(\{\phi_n, t_n\} | \pi_k) = \ln \Pi_{n=1}^N \left[\pi_1 \mathbb{P}(\phi | \mathcal{C}_1)^{t_n} \dots \pi_K \mathbb{P}(\phi | \mathcal{C}_K)^{t_n} \right]$$

$$\ln \mathbb{P}(\{\phi_n, t_n\} | \pi_k) = \ln \Pi_{n=1}^N \Pi_{k=1}^K \left[\pi_k \mathbb{P}(\phi | \mathcal{C}_k) \right]^{t_n}$$

$$\ln \mathbb{P}(\{\phi_n, t_n\} | \pi_k) = \sum_{n=1}^N \sum_{k=1}^K \ln \left(\left[\pi_k \mathbb{P}(\phi | \mathcal{C}_k) \right]^{t_n} \right)$$

$$\ln \mathbb{P}(\{\phi_n, t_n\} | \pi_k) = \sum_{n=1}^N \sum_{k=1}^K t_n \left[\ln \pi_k + \ln \mathbb{P}(\phi | \mathcal{C}_k) \right]$$

2 -

we will denote t_{nk} the k-th element of the vector \mathbf{t}_n . We have the following likelihood and log likelihood:

$$\mathcal{L}(t_n | \pi_k) = \Pi_{n=1}^N \Pi_{k=1}^K \pi_k^{t_{nk}}$$

$$l(t_n | \pi_k) = \sum_{n=1}^N \sum_{k=1}^K t_{nk} \ln(\pi_k)$$

We are looking at the following optimization problem: $\begin{cases} \max l(t_n | \pi_k) \\ \sum \pi_k - 1 = 0 \end{cases}$

which is equivalent to this one: $\max_{\pi_k} L = l(t_n | \pi_k) + \lambda \left[\sum_{i=1}^K \pi_i \right]$

$$\iff \max_{\pi_k} L = \sum_{n=1}^N \sum_{i=1}^K t_{nk} \ln(\pi_i) + \lambda \left[\sum_{i=1}^K \pi_i \right]$$

$$\frac{\partial L}{\partial \pi_k} = 0 \iff \sum_{n=1}^N t_{nk} \frac{1}{\pi_k} + \lambda = 0 \iff \pi_k = \frac{1}{\lambda} \sum_{n=1}^N t_{nk}$$

$$\text{But we also have } \frac{\partial L}{\partial \lambda} = 0 \implies \sum_{i=1}^K \pi_i - 1 = 0$$

By mixing both equation we get:

$$\frac{-1}{\lambda} \sum_{i=1}^K \sum_{n=1}^N t_{nk} = 1 \implies \frac{-N}{\lambda} = 1 \implies \lambda = -N$$

We finally get:

$$\hat{\pi}_k = \frac{1}{N} \sum_{n=1}^N t_{nk}$$

which is intuitive. The estimator is the proportion of number elements in each class.

3 -

$$l = \ln \mathbb{P}(\{\phi_n, t_n\} | \pi_k) = \sum_{n=1}^N \sum_{k=1}^K t_n \left[\ln \pi_k + \ln \mathbb{P}(\phi | \mathcal{C}_k) \right] = \sum_{n=1}^N \sum_{k=1}^K t_n \left[\ln \pi_k + \ln \mathcal{N}(\phi | \mu_k, \Sigma) \right]$$

$$l = \sum_{n=1}^N \sum_{k=1}^K t_n \left[\ln \pi_k + \ln |\Sigma|^{-1/2} + \ln \left(\frac{1}{\sqrt{2\pi}} \right) - \frac{1}{2} [(x_n - \mu_k)^T \Sigma^{-1} (x_n - \mu_k)] \right]$$

$$\nabla_{\mu_k} l = \sum_{n=1}^N \nabla_{\mu_k} - \frac{t_{nk}}{2} [(x_n - \mu_k)^T \Sigma^{-1} (x_n - \mu_k)]$$

4-

$$\nabla_{\mu_k} l = 0 \iff \sum_{n=1}^N t_{nk}(x_n - \mu_k) = 0$$

we will denote $\sum_{n=1}^N t_{nk} = N_k$

$$\hat{\mu}_k = \frac{1}{N_k} \sum_{n=1}^N t_{nk} x_n$$

which is the mean of all the input vector x_n assigned to \mathcal{C}_k

5-

$$\nabla_{\Sigma} l = \sum_{n=1}^N \sum_{k=1}^K \left[t_{nk} \nabla_{\Sigma} \ln |\Sigma|^{-1/2} - \frac{t_{nk}}{2} \nabla_{\Sigma} (x_n - \mu_k)^T \Sigma^{-1} (x_n - \mu_k) \right]$$

$$\nabla_{\Sigma} l = \sum_{n=1}^N \sum_{k=1}^K \left[\frac{-t_{nk}}{2} \nabla_{\Sigma} \ln |\Sigma| - \frac{t_{nk}}{2} \nabla_{\Sigma} (x_n - \mu_k)^T \Sigma^{-1} (x_n - \mu_k) \right]$$

Thanks to the cookbook formulas, we get:

$$\nabla_{\Sigma} l = \sum_{n=1}^N \sum_{k=1}^K \frac{-t_{nk}}{2} \Sigma^{-1} + \sum_{n=1}^N \sum_{k=1}^K \frac{t_{nk}}{2} (\Sigma^{-T} (x_n - \mu_k)(x_n - \mu_k)^T \Sigma^{-T})$$

Σ is a covariance matrix hence is symmetrical by construction, we get $\Sigma^T = \Sigma$

$$\nabla_{\Sigma} l = -\Sigma^{-1} \frac{N}{2} + \sum_{k=1}^K \sum_{n \in \mathcal{C}_k} \left[\frac{1}{2} \Sigma^{-1} (x_n - \mu_k)(x_n - \mu_k)^T \Sigma^{-1} \right]$$

$$\nabla_{\Sigma} l = -\Sigma^{-1} \frac{N}{2} + \frac{1}{2} \Sigma^{-1} \sum_{k=1}^K \sum_{n \in \mathcal{C}_k} \left[(x_n - \mu_k)(x_n - \mu_k)^T \right] \Sigma^{-1}$$

6 -

$$\nabla_{\Sigma} l = 0 \iff N = \Sigma^{-1} \sum_{k=1}^K \sum_{n \in \mathcal{C}_k} (x_n - \mu_k)(x_n - \mu_k)^T$$

$$\hat{\Sigma} = \frac{1}{N} \sum_{k=1}^K \sum_{n \in \mathcal{C}_k} (x_n - \mu_k)(x_n - \mu_k)^T$$

We will denote $S_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} (x_n - \mu_k)(x_n - \mu_k)^T$ the covariance associated with each class.

$$\hat{\Sigma} = \sum_{k=1}^K \frac{N_k}{N} S_k$$

Again, this is intuitive, it's the weighted average of the covariance matrices associated with each class.

3. Classifying Fruit [15pts]

You're tasked with classifying three different kinds of fruit, based on their heights and widths. Figure 1 is a plot of the data. Iain Murray collected these data and you can read more about this on his website at http://homepages.inf.ed.ac.uk/imurray2/teaching/oranges_and_lemons/. We have made a slightly simplified (collapsing the subcategories together) version of this available as `fruit.csv`, which you will find in the Github repository. The file has three columns: `type` (1=apple, 2=orange, 3=lemon), `width`, and `height`. The first few lines look like this:

```
fruit,width,height
1,8.4,7.3
1,8,6.8
1,7.4,7.2
1,7.1,7.8
...
```




Figure 1: Heights and widths of apples, oranges, and lemons. These fruit were purchased and measured by Iain Murray: http://homepages.inf.ed.ac.uk/imurray2/teaching/oranges_and_lemons/.

Problem 3 (Classifying Fruit, 15pts)

Please implement the following:

- Implement the three-class generalization of logistic regression, also known as softmax regression, for these data. You will do this by implementing gradient descent on the log likelihood.
- After this, implement a simple generative classifier with Gaussian class-conditional densities, as in Bishop Section 4.2.2. In particular, make two implementations of this, one with a shared covariance matrix across all of the classes, and one with a separate covariance being learned for each class. Note that the staff implementation can switch between these two by the addition of just a few lines of code. The shared covariance matrix case is detailed in Bishop (and you worked on it in Problem 2), and the separate covariance case is only slightly different. In the separate covariance matrix case, the MLE for the covariance matrix of each class is simply the covariance of the data points assigned to that class, without combining them as in the shared case.

You may use anything in numpy or scipy, except for scipy.optimize. That being said, if you happen to find a function in numpy or scipy that seems like it is doing too much for you, run it by a staff member. In general, linear algebra and random variable functions are fine. The controller file is problem3.py, in which you will specify parameters. The actual implementations you will write will be in LogisticRegression.py and GaussianGenerativeModel.py.

You will be given unimplemented class interfaces for GaussianGenerativeModel and LogisticRegression in the distribution code, and the code will indicate certain lines that you should not change in your final submission. Naturally, don't change these. These classes will allow the final submissions to have consistency. There will also be a few hyperparameters that are set to irrelevant values at the moment. You may need to modify these to get your methods to work. The classes you implement follow the same pattern as scikit-learn, so they should be familiar to you. The distribution code currently outputs nonsense predictions just to show what the high-level interface should be, so you should completely remove the given predict() implementations and replace them with your implementations.

- The visualize() method for each classifier will save a plot that will show the decision boundaries. Please include those in this assignment.
- Which classifiers model the distributions well?
- What explains the differences?

Solution

We started with a Logistic classification implemented with a Gradient descent algorithm. We used the gradient's formula from the 5th question of problem 1: $\nabla_{w_j} E(w_1, \dots, w_K) = \sum_{n=1}^N \phi_n(y_{nj} - t_{nj})$

We then used the following model for the Generative Gaussian with shared variance:

$$p(C_k|x) = \frac{\exp(a_k)}{\sum_j \exp(a_j)} \text{ with } a_k(x) = \mathbf{w}_k^T x + w_{k0}, \mathbf{w}_k = \Sigma^{-1} \mu_k \text{ and } w_{k0} = -\frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \ln p(C_k)$$

We finally looked at the case with different variances with the following model:

$$p(C_k|x) = \frac{\exp(a_k)}{\sum_j \exp(a_j)} \text{ with } a_k(x) = x^T \mathbf{W}_{k2}^T x + \mathbf{w}_{k1}^T x + w_{k0}^1$$

$$\mathbf{W}_{k2} = -\frac{1}{2} \Sigma_k^{-1}, \mathbf{w}_{k1} = \Sigma_k^{-1} \mu_k \text{ and } w_{k0} = -\frac{1}{2} \mu_k^T \Sigma_k^{-1} \mu_k - \frac{1}{2} \ln |\Sigma_k| + \ln P(C_k)$$

The Generative Gaussian with separate variances fits the best thanks to the quadratic equation in x .

¹ $x^T \mathbf{W}_{k2}^T x$ is the Mahalanobis distance




Figure 2: Logistic Regression with Gradient Descent

Calibration [1pt]

Approximately how long did this homework take you to complete? 10h




Figure 3: Generative Gaussian with Shared Covariances




Figure 4: Generative Gaussian with Separated Covariances